

# ExecutionRank: Verified Multi-Agent Execution via Weighted Attestation and Threshold Acceptance

Nikko Ambroselli  
daybed-wile-0e@icloud.com

February 14, 2026

## Abstract

Autonomous agent systems increasingly delegate tool invocation to remote agents discovered through registries. While this enables scale and specialization, it introduces a missing primitive: *verifiable execution trust*. Most systems accept a tool’s return value as evidence of correctness or rely on an LLM to judge outputs, neither of which constitutes independent verification. We present *ExecutionRank*, a trust layer for autonomous tool invocation that combines (i) verified multi-agent execution (worker + independent verifiers), (ii) a weighted attestation graph built from auditable positive attestations, and (iii) deterministic threshold acceptance to gate consumption of results (and optionally, authorization of side effects). ExecutionRank yields a policy-tunable assurance signal for tool calls under uncertainty.

**Keywords:** autonomous agents; tool invocation; verification; attestation; reputation; EigenTrust; risk gating; audit receipts.

## 1 Introduction and Problem Context

Autonomous agent ecosystems delegate tool invocation to remote agents discovered via registries (e.g., ERC-8004). A host ( $H$ ) selects an agent, invokes a tool, and consumes the result. This model lacks verifiable execution trust.

**Unverified execution.** Most systems treat “tool returned successfully” as evidence of correctness. A faulty or malicious worker can return arbitrary data.

**LLM-as-judge insufficiency.** Semantic plausibility checks are subjective and vulnerable to prompt injection and distribution shift.

**Absence of independent attestation.** Traditional distributed systems introduce independent witnesses (replication, quorum, BFT [1, 4]). Agent tool invocation typically lacks such witnesses.

**Risk of unverified invocation.** Tools may have side effects (payments, state mutation, external API calls). Accepting unverified results risks incorrect state and financial loss.

**Contributions.** ExecutionRank introduces (1) a verified execution protocol, (2) a weighted attestation graph, and (3) deterministic threshold acceptance.

## 2 System Model

### 2.1 Entities and Roles

We consider a host  $H$ , a control plane  $CP$ , agents exposing tools, and a receipt store  $RS$ .

Role	Identity	Responsibility
Worker ( $W$ )	Agent ID	Executes tool and produces outputs
Verifier ( $V$ )	Agent ID	Independently validates outputs
Invoker ( $I$ )	Principal	Consumes result; may attest semantics (post-hoc via a

Table 1: Roles in verified multi-agent execution.

### 2.2 Assumptions and Threat Model (Outline)

We assume  $RS$ s append-only (or tamper-evident) and that verifiers can be sampled from a pool not fully controlled by a single worker. We consider adversaries controlling subsets of workers, verifiers, or invokers (Section 7).

## 3 Protocol Overview

### 3.1 Verified Multi-Agent Execution

Given a request to invoke tool  $T$  on worker  $W$  with args  $x$ : (i)  $CP$  invokes  $T$  on  $W$  to obtain  $y$ , (ii) computes digests  $d_x \leftarrow H(x)$  and  $d_y \leftarrow H(y)$ , (iii) samples  $k$  verifiers  $V_1..V_k$  with  $V_i \neq W$ , (iv) collects verification receipts

(verdict, confidence, mode), (v) stores receipts, and (vi) decides acceptance (Section 6).

## 4 Verification Modes

Modes are ordered by strength:

- REPLAY: re-execute and compare outputs (or validate invariants for non-deterministic tools).
- CROSSCHECK: validate via alternate logic or independent sources.
- HEURISTIC: rule-based or approximate checks.

## 5 Weighted Attestation Graph

### 5.1 Graph Structure and Rationale

We construct a bipartite graph with edges from attestors to workers; only PASSattestations create edges. Worker→worker edges are disallowed to prevent mutual inflation loops.

### 5.2 Edge Weight

For attester  $A$  attesting PASS to worker  $W$ , define:

$$t(A, W) = \lambda_{\text{role}} \cdot \text{conf}(A, W) \cdot s(\text{mode}) \cdot g(\text{agree}) \cdot \delta(\Delta t). \quad (1)$$

Implementation:  $\lambda_{\text{verifier}} = 1$ ,  $\lambda_{\text{invoker}} = 0.5$ ;  $s(\text{REPLAY}) = 1$ ,  $s(\text{CROSSCHECK}) = 0.8$ ,  $s(\text{HEURISTIC}) = 0.6$ ;  $g$  boosts when multiple attestors concur on the same task (capped 1.2);  $\delta = \exp(-\ln 2 \cdot \Delta t / \text{halfLife})$  with half-life 90 days.

### 5.3 Normalized Matrix $C$

$$\mathbf{C}_{A,W} = \frac{\max(t(A, W), 0)}{\sum_{W'} \max(t(A, W'), 0)}. \quad (2)$$

#### 5.3.1 Attester–Attester Flow Matrix $M$ (Derived from $C$ )

Let  $\mathcal{A}$  be the set of attestors ( $|\mathcal{A}| = n$ ) and  $\mathcal{W}$  the set of workers ( $|\mathcal{W}| = m$ ). Let  $C \in \mathbb{R}^{n \times m}$  be the row-stochastic attester→worker matrix defined in Section 5.3.

Define the worker column-sums:

$$s_w = \sum_{a \in \mathcal{A}} C_{a,w} \quad \forall w \in \mathcal{W}.$$

Define a worker→attester redistribution matrix  $Q \in \mathbb{R}^{m \times n}$ :

$$Q_{w,b} = \begin{cases} \frac{C_{b,w}}{s_w}, & s_w > 0, \\ \frac{1}{n}, & s_w = 0, \end{cases} \quad \forall w \in \mathcal{W}, b \in \mathcal{A}.$$

Then the induced attester→attester flow matrix is:

$$\begin{aligned} M &= CQ \\ &= CD^{-1}C^\top. \end{aligned}$$

where  $D = \text{diag}(s_1, \dots, s_m)$  (and  $D^{-1}$  is interpreted with the  $s_w = 0$  convention above).

Equivalently, entrywise:

$$M_{a,b} = \sum_{w \in \mathcal{W}} C_{a,w} \frac{C_{b,w}}{\sum_{a' \in \mathcal{A}} C_{a',w}} \quad (\text{with } 0/0 \text{ handled by the } s_w = 0 \text{ case}).$$

**Row-stochasticity.** If  $C$  is row-stochastic and  $Q$  is defined as above, then  $M$  is row-stochastic:  $\sum_b M_{a,b} = 1$  for all  $a$ .

### 5.4 Seed Vector $\mathbf{p}$ and Baseline Rank

The seed (teleport) vector  $\mathbf{p}$  determines the reputation mass each attester receives when not following the attestation graph. To ensure every attester—including those with no attestation history—receives a well-defined baseline, we define

$$\mathbf{p} = (1 - \beta) \mathbf{u} + \beta \mathbf{s}, \quad (3)$$

where  $\mathbf{u}$  is the uniform distribution over  $\mathcal{A}$  (all attestors who have appeared in any receipt),  $\mathbf{s}$  is a seeded distribution (e.g., prior trust or domain-specific weights), and  $\beta \in [0, 1]$ . The uniform component guarantees  $\min_{a \in \mathcal{A}} p_a \geq (1 - \beta)/|\mathcal{A}|$ ; we denote this baseline  $r_0$ . Choosing  $\beta$  (e.g.,  $\beta = 0.8$ ) yields a tunable  $r_0$ ; default  $r_0 \approx 0.2$  for small  $|\mathcal{A}|$ .

### 5.5 Reputation via EigenTrust-Style Iteration

Following the EigenTrust approach [3]:

$$\mathbf{R} = (1 - \alpha)\mathbf{p} + \alpha \mathbf{M}\mathbf{R}. \quad (4)$$

Attestors with no attestation history (rows of zeros in  $\mathbf{C}$ ) receive reputation solely from the teleport term; their rank is at least  $r_0$  from  $\mathbf{p}$ .

## 6 Acceptance Threshold Logic

Let  $\mathcal{A}$  be PASS attestors for a task, and  $R_A$  their reputations.

**Definition 1** (Acceptance Predicate).

$$\text{accepted} \iff \left( \sum_{A \in \mathcal{A}} R_A \geq \tau \right) \wedge (|\mathcal{A}| \geq 1). \quad (5)$$

$\mathcal{A}$  is the set of verifier attestors that returned PASS for the task (invoker attestation is asynchronous and does not affect immediate acceptance). Attestors not yet in the attestation graph receive  $R_A = r_0$  from the seed vector  $\mathbf{p}$  (Section 5.4). Default  $\tau = 0.5$ ,  $\alpha = 0.85$ .

---

**Algorithm 1** VerifiedExecution (Control Plane)

**Require:** Worker  $W$ , tool  $T$ , args  $x$ , verifier count  $k$ , policy  $\Pi$   
**Ensure:** Decision accepted/rejected, receipts

```
1:  $y \leftarrow \text{CP.Invoke}(W, T, x)$ 
2:  $d_x \leftarrow H(x); d_y \leftarrow H(y)$ 
3:  $\mathcal{V} \leftarrow \text{CP.SelectVerifiers}(k, \Pi, \text{exclude} = W)$ 
4: for all  $V \in \mathcal{V}$  do
5:    $(v, c, m) \leftarrow \text{CP.Invoke}(V, \text{verify\_result}, (\text{taskId}, W, T, d_x, d_y))$             $\triangleright d_x = H(x), d_y = H(y); \text{ verifier returns}$ 
     {verdict, confidence, mode}
6:    $\text{CP.AppendVerificationReceipt}(V, W, T, d_x, d_y, v, c, m)$ 
7: end for
8:  $\text{CP.AppendExecutionReceipt}(W, T, d_x, d_y)$ 
9: return  $\text{CP.DecideAcceptance}(\text{taskId})$ 
```

---

**Algorithm 2** ComputeAttectorReputation

**Require:** PASS receipts, damping  $\alpha$ , seed  $\mathbf{p}$ , tolerance  $\varepsilon$ , maxIters  
**Ensure:** Attector reputation vector  $\mathbf{R}$

```
1: Build  $t(A, W)$  from PASS receipts
2: Normalize rows to form  $\mathbf{C}$ 
3: Derive  $\mathbf{M}$  from  $\mathbf{C}$  (attector agreement flow through shared workers)
4:  $\mathbf{R} \leftarrow \mathbf{p}$ 
5: for  $i = 1$  to maxIters do
6:    $\mathbf{R}_{\text{next}} \leftarrow (1 - \alpha)\mathbf{p} + \alpha(\mathbf{M}\mathbf{R})$ 
7:   if  $\|\mathbf{R}_{\text{next}} - \mathbf{R}\|_1 < \varepsilon$  then break
8:   end if
9:    $\mathbf{R} \leftarrow \mathbf{R}_{\text{next}}$ 
10: end for
11: Normalize  $\mathbf{R}$  (e.g.,  $\max(\mathbf{R}) = 1$ )
12: return  $\mathbf{R}$ 
```

---

## 7 Security and Adversarial Considerations

### 7.1 Sybil Attacks

Sybil attacks [2] allow a single entity to control multiple identities. Attack / mitigation / residual risk.

### 7.2 Collusion Rings

Attack / mitigation / residual risk.

### 7.3 Low-Risk Farming

Attack / mitigation / residual risk.

### 7.4 Invoker-Only Inflation

Attack / mitigation / residual risk.

### 7.5 Receipt Integrity and Omission

Attack / mitigation / residual risk.

## 8 Deployment and Integration Models

Centralized control plane; decentralized receipt sets; enterprise tool invocation; payment-gated execution.

## 9 Implementation Notes

The attest-substrate control plane implements Execution-Rank: discovery via ERC-8004 subgraph (search\_agents, list\_tools); policy filters (MIN\_REPUTATION, ALLOW\_RISK\_CLASS); invoke always runs worker + verifiers; receipts stored in SQLite; ranking refreshed periodically (default 60s). Verifier selection: allowlist (capped at 10 agents) or search-based; worker excluded. Verifier tool: verify\_result (fallback validate); input {taskId, workerAgentId, workerTool, workerArgsHash, workerOutputHash}; output {verdict, confidence, mode}. Digests: SHA-256 of canonical JSON.

## 10 Evaluation Plan

Fault injection; lazy verifiers; collusion simulation; latency/cost vs assurance; convergence under sparsity.

## 11 Limitations and Future Work

Negative evidence; risk-class reputation; verifier diversity guarantees; receipt consensus; privacy; incentives.

## A Notation

Symbol	Meaning
H	Host
CP	Control plane
RS	Receipt store
W	Worker agent
V	Verifier agent
A	Attestor
C	Attestor→worker matrix
M	Attestor→attestor flow matrix
R	Attestor reputation vector
p	Seed distribution $(1 - \beta)\mathbf{u} + \beta\mathbf{s}$
$r_0$	Baseline rank (from uniform component of p)
$\alpha$	Damping
$\beta$	Seed mixing (uniform vs seeded)
$\tau$	Acceptance threshold

Table 2: Notation summary.

## B Illustrative Receipt Schemas

### B.1 Execution Receipt (Example)

## References

- [1] Miguel Castro and Barbara Liskov. Practical byzantine fault tolerance. In *Proceedings of the Third USENIX Symposium on Operating Systems Design and Implementation (OSDI), New Orleans, Louisiana, USA, February 22–25, 1999*, pages 173–186. USENIX Association, 1999.
- [2] John R. Douceur. The sybil attack. In Peter Druschel, Frans Kaashoek, and Antony Rowstron, editors, *Peer-to-Peer Systems: First International Workshop, IPTPS 2002, Cambridge, MA, USA, March 7–8, 2002, Revised Papers*, volume 2429 of *Lecture Notes in Computer Science*, pages 251–260. Springer, 2002.
- [3] Sepandar D. Kamvar, Mario T. Schlosser, and Hector Garcia-Molina. The eigentrust algorithm for reputation management in P2P networks. In *Proceedings of the Twelfth International World Wide Web Conference, WWW 2003, Budapest, Hungary, May 20–24, 2003*, pages 640–651. ACM, 2003.
- [4] Leslie Lamport, Robert E. Shostak, and Marshall C. Pease. The byzantine generals problem. *ACM Transactions on Programming Languages and Systems*, 4(3):382–401, 1982.

```
{  
    "receipt_id": "uuid",  
    "task_id": "uuid",  
    "principal": "anonymous",  
    "agent_id": 12345,  
    "tool_name": "coinbase_price",  
    "tool_id": "12345::coinbase_price",  
    "args_digest": "sha256-hex...",  
    "result_digest": "sha256-hex...",  
    "outcome": "ok",  
    "started_at": "ISO8601",  
    "ended_at": "ISO8601",  
    "latency_ms": 42  
}
```

Figure 1: Execution Receipt (attest-substrate schema). Digests are SHA-256 of canonical JSON.